

Traffic Sign Recognition with Data Augmentation

Overview

This project implements a **deep learning pipeline** for the recognition of **German Traffic Signs** using **Convolutional Neural Networks (CNNs)**.

The goal is to **accurately classify traffic signs into 43 categories**, a crucial step in developing **autonomous driving systems** and **advanced driver-assistance systems (ADAS)**.

Traffic sign recognition plays a critical role in **road safety**, enabling vehicles to automatically detect and interpret signs such as **speed limits, stop signs, and warnings**.

Objectives

- Build a **robust classification model** for the German Traffic Sign Recognition Benchmark (GTSRB).
- Apply **data augmentation** to improve generalization.
- Compare baseline and augmented results.
- Evaluate model with **accuracy, precision, recall, and F1-score**.
- Provide a **reproducible training pipeline** for researchers and engineers.

Dataset

- **Name:** German Traffic Sign Recognition Benchmark (GTSRB)
- **Number of classes:** 43
- **Training images:** ~39,000
- **Test images:** ~12,000
- **Image size:** Resized to 32×32×3

Example Images

Original Augmented

Class Distribution

The dataset is **imbalanced**, with some classes having thousands of samples (e.g., speed limits) while others have fewer (e.g., rare signs).

Methodology

1. Data Preprocessing

- **Resizing** all images to 32×32×3.
- **Normalization** to range [0, 1].
- **One-hot encoding** of labels.

2. Data Augmentation

To improve generalization, the following transformations were applied:

- Random rotations ($\pm 20^\circ$)
- Zoom (0.8–1.2×)
- Horizontal/vertical shifts ($\pm 10\%$)
- Brightness variation ($\pm 20\%$)
- Shear transformations

- Horizontal flips

3. Model Architecture (CNN)

Input: 32x32x3

↓

Conv2D (32 filters, 3x3) + ReLU

↓

Conv2D (64 filters, 3x3) + ReLU

↓

MaxPooling (2x2)

↓

Dropout (0.25)

↓

Conv2D (128 filters, 3x3) + ReLU

↓

MaxPooling (2x2)

↓

Dropout (0.25)

↓

Flatten

↓

Dense (256 units) + ReLU

↓

Dropout (0.5)

↓

Dense (43 units) + Softmax

Training Setup

- **Loss Function:** Categorical Crossentropy
- **Optimizer:** Adam (lr=0.001)
- **Batch Size:** 32 / 64
- **Epochs:** 30–50
- **Callbacks:**
 - EarlyStopping (patience=5)

- ModelCheckpoint (best model saved)

Results

Accuracy

- **Baseline CNN:** ~92% test accuracy
- **With Augmentation:** 95–97% test accuracy

Evaluation Metrics

| Metric | Score |
|-----------|-------|
| Accuracy | 95% |
| Precision | 94% |
| Recall | 94% |
| F1-score | 94% |

Confusion Matrix

Training Curves

Reproducibility

1. Clone Repo

```
git clone https://github.com/your-username/traffic-signs-classification.git  
cd traffic-signs-classification
```

2. Install Dependencies

```
pip install -r requirements.txt
```

3. Run Training Notebook

```
jupyter notebook augmented_traffic_signs.ipynb
```

4. Evaluate Model

```
from tensorflow.keras.models import load_model  
model = load_model("best_model.h5")
```

Dependencies

- Python 3.8+
- TensorFlow / Keras
- NumPy
- Matplotlib
- OpenCV
- Scikit-learn

Install with:

```
pip install tensorflow keras numpy matplotlib opencv-python scikit-learn
```

Future Work

- Apply **transfer learning** with ResNet, VGG, or EfficientNet.
- Experiment with **semi-supervised learning** for rare signs.
- Deploy model using **Flask, FastAPI, or Streamlit**.

- Optimize for **edge devices** (Jetson Nano, Raspberry Pi).

Acknowledgements

- GTSRB Dataset
- TensorFlow & Keras documentation
- OpenCV tutorials